

Importing Libraries

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import re
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error
from sklearn.datasets import load_boston
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

Problem Statement

- You have been given a dataset that describes the houses in Boston. Now, based on the given features, you have to predict the house price.

Creating a DataFrame

```
boston = load_boston()
df = pd.DataFrame(boston.data)
```

EDA - Exploratory Data Analysis

```
df.head()
```

	0	1	2	3	4	5	6	7	8	9	10	11	12
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33

- Adding columns

```
df.columns = boston.feature_names
```

```
df.head()
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LS
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33

Columns Informations

- CRIM per capita crime rate by town
- ZN proportion of residential land zoned for lots over 25,000 sq.ft.
- INDUS proportion of non-retail business acres per town
- CHAS Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
- NOX nitric oxides concentration (parts per 10 million)
- RM average number of rooms per dwelling
- AGE proportion of owner-occupied units built prior to 1940
- DIS weighted distances to five Boston employment centres

- RAD index of accessibility to radial highways
- TAX full-value property-tax rate per 10,000usd
- PTRATIO pupil-teacher ratio by town
- B 1000(Bk - 0.63)^2 where Bk is the proportion of blacks by town
- LSTAT % lower status of the population

✓ Adding the target column into the DataFrame

```
df['PRICE'] = boston.target
```

```
df.head()
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LS
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	

```
df.tail()
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	L
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1.0	273.0	21.0	391.99	
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1.0	273.0	21.0	396.90	
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1.0	273.0	21.0	396.90	
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1.0	273.0	21.0	393.45	
505	0.04741	0.0	11.93	0.0	0.573	6.030	80.8	2.5050	1.0	273.0	21.0	396.90	

```
df.shape
```

(506, 14)

```
df.columns
```

Index(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX',  
 'PTRATIO', 'B', 'LSTAT', 'PRICE'],  
 dtype='object')

```
df.dtypes
```

CRIM float64  
ZN float64  
INDUS float64  
CHAS float64  
NOX float64  
RM float64  
AGE float64  
DIS float64  
RAD float64  
TAX float64  
PTRATIO float64  
B float64  
LSTAT float64  
PRICE float64  
dtype: object

```
df.nunique()
```

CRIM 504  
ZN 26  
INDUS 76  
CHAS 2  
NOX 81  
RM 446  
AGE 356  
DIS 412  
RAD 9  
TAX 66  
PTRATIO 46

```
B          357
LSTAT      455
PRICE      229
dtype: int64

df.isnull()

   CRIM  ZN  INDUS  CHAS  NOX  RM  AGE  DIS  RAD  TAX  PTRATIO  B  L
0  False False False False False False False False False False False False
1  False False False False False False False False False False False False
2  False False False False False False False False False False False False
3  False False False False False False False False False False False False
4  False False False False False False False False False False False False
...    ...  ...    ...    ...    ...    ...    ...    ...    ...    ...    ...    ...
501 False False False False False False False False False False False False
502 False False False False False False False False False False False False
503 False False False False False False False False False False False False
504 False False False False False False False False False False False False
505 False False False False False False False False False False False False

506 rows x 14 columns

df.isnull().sum()

CRIM      0
ZN        0
INDUS     0
CHAS      0
NOX       0
RM        0
AGE       0
DIS       0
RAD       0
TAX       0
PTRATIO   0
B         0
LSTAT     0
PRICE     0
dtype: int64

df.describe()

   CRIM  ZN  INDUS  CHAS  NOX  RM  AG
count  506.000000  506.000000  506.000000  506.000000  506.000000  506.000000  506.000000
mean    3.613524   11.363636   11.136779    0.069170    0.554695    6.284634   68.57490
std     8.601545   23.322453    6.860353    0.253994    0.115878    0.702617   28.14886
min     0.006320    0.000000    0.460000    0.000000    0.385000    3.561000    2.90000
25%     0.082045    0.000000    5.190000    0.000000    0.449000    5.885500   45.02500
50%     0.256510    0.000000    9.690000    0.000000    0.538000    6.208500   77.50000
75%     3.677083   12.500000   18.100000    0.000000    0.624000    6.623500   94.07500
max     88.976200  100.000000   27.740000    1.000000    0.871000    8.780000  100.00000

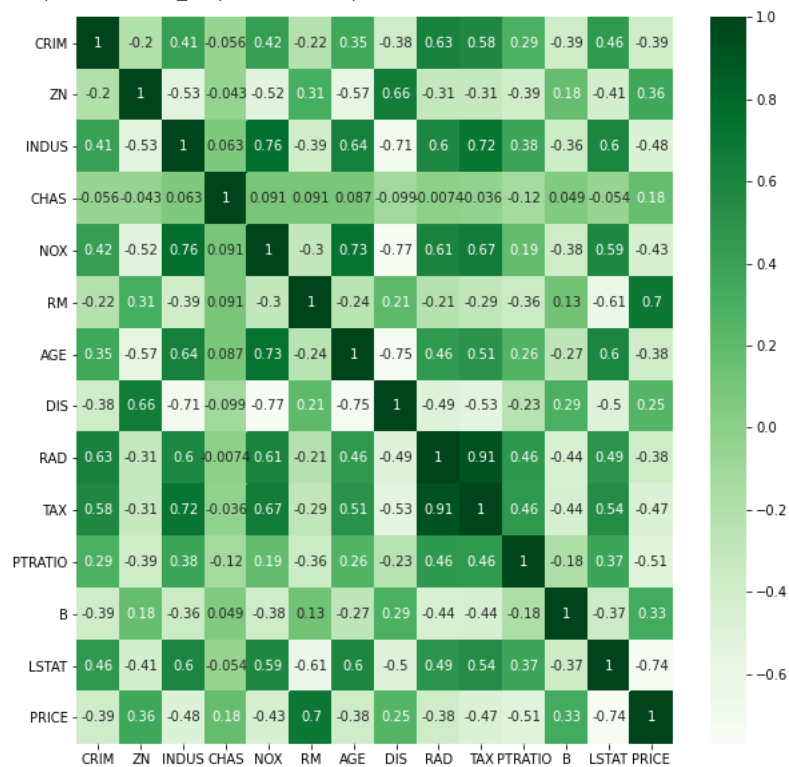
df.corr

<bound method DataFrame.corr of
0  0.00632  18.0  2.31  0.0  0.538  ...  296.0  15.3  396.90  4.98  24.0
1  0.02731  0.0  7.07  0.0  0.469  ...  242.0  17.8  396.90  9.14  21.6
2  0.02729  0.0  7.07  0.0  0.469  ...  242.0  17.8  392.83  4.03  34.7
3  0.03237  0.0  2.18  0.0  0.458  ...  222.0  18.7  394.63  2.94  33.4
4  0.06905  0.0  2.18  0.0  0.458  ...  222.0  18.7  396.90  5.33  36.2
..  ...  ...  ...  ...  ...  ...  ...  ...  ...  ...  ...
501 0.06263  0.0  11.93  0.0  0.573  ...  273.0  21.0  391.99  9.67  22.4
502 0.04527  0.0  11.93  0.0  0.573  ...  273.0  21.0  396.90  9.08  20.6
503 0.06076  0.0  11.93  0.0  0.573  ...  273.0  21.0  396.90  5.64  23.9
504 0.10959  0.0  11.93  0.0  0.573  ...  273.0  21.0  393.45  6.48  22.0
505 0.04741  0.0  11.93  0.0  0.573  ...  273.0  21.0  396.90  7.88  11.9

[506 rows x 14 columns]>
```

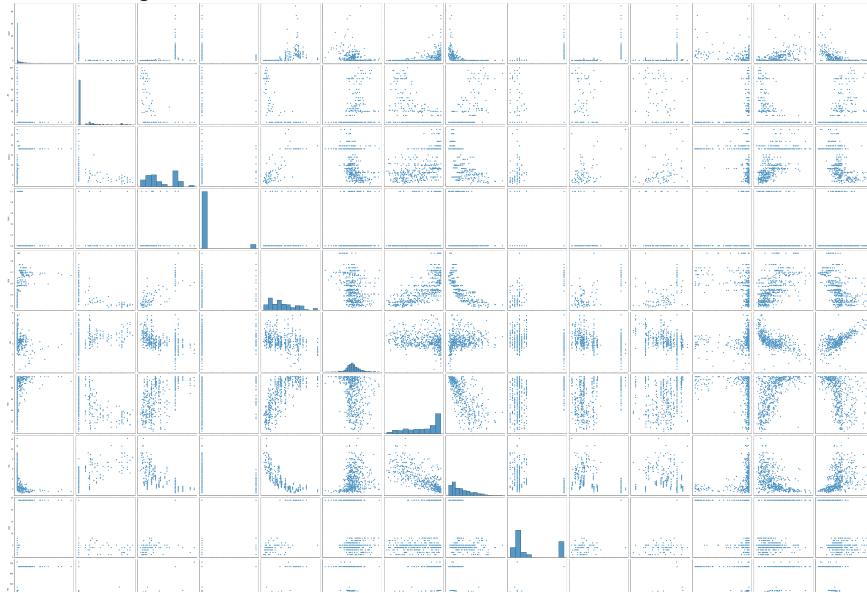
```
plt.figure(figsize=(10,10))
sns.heatmap(data=df.corr(), annot=True, cmap='Greens')
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7fe05ed6ff50>



```
sns.pairplot(df, size=5)
```


```
/usr/local/lib/python3.7/dist-packages/seaborn/axisgrid.py:1969: UserWarning: The `si
warnings.warn(msg, UserWarning)
<seaborn.axisgrid.PairGrid at 0x7fe057603890>
```



```
# Plot a Boxplot
plt.figure(figsize=(50,50))
df.boxplot()
```



<matplotlib.axes.\_subplots.AxesSubplot at 0x7fe04b646950>



```
# Minimum Price
df.PRICE.min()
```

5.0

```
# Maximum Price
df.PRICE.max()
```

50.0

```
# Standard Deviation
df.PRICE.std()
```

9.19710408737982

## Export the dataset

```
df.to_csv('boston_dataset.csv',)
```

## Machine Learning - Linear Regression

```
df.head()
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LS
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	✓
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	!
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	✓
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	;
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	!

```
X = np.array(df.drop('PRICE', axis=1))
y = np.array(df.PRICE)
```

```
# X = boston.data
# y = boston.target
```

## Splitting the data

```
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.2,
                                                    random_state=42)
```

```
len(X_train)
```

404

```
len(y_train)
```

404

```
len(X_test)
```

```
102
```

```
len(y_test)
```

```
102
```

## ✓ Choosing the model

```
model = LinearRegression()
```

## ✓ Fitting/Train the model

```
model.fit(X_train,y_train)
```

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
# Intercept Value
```

```
model.intercept_
```

```
30.24675099392392
```

```
# Coefficient Value
```

```
model.coef_
```

```
array([-1.13055924e-01,  3.01104641e-02,  4.03807204e-02,  2.78443820e+00,
        -1.72026334e+01,  4.43883520e+00, -6.29636221e-03, -1.44786537e+00,
         2.62429736e-01, -1.06467863e-02, -9.15456240e-01,  1.23513347e-02,
        -5.08571424e-01])
```

## ✓ Prediction

```
y_test
```

```
array([23.6, 32.4, 13.6, 22.8, 16.1, 20. , 17.8, 14. , 19.6, 16.8, 21.5,
        18.9,  7. , 21.2, 18.5, 29.8, 18.8, 10.2, 50. , 14.1, 25.2, 29.1,
        12.7, 22.4, 14.2, 13.8, 20.3, 14.9, 21.7, 18.3, 23.1, 23.8, 15. ,
        20.8, 19.1, 19.4, 34.7, 19.5, 24.4, 23.4, 19.7, 28.2, 50. , 17.4,
        22.6, 15.1, 13.1, 24.2, 19.9, 24. , 18.9, 35.4, 15.2, 26.5, 43.5,
        21.2, 18.4, 28.5, 23.9, 18.5, 25. , 35.4, 31.5, 20.2, 24.1, 20. ,
        13.1, 24.8, 30.8, 12.7, 20. , 23.7, 10.8, 20.6, 20.8,  5. , 20.1,
        48.5, 10.9,  7. , 20.9, 17.2, 20.9,  9.7, 19.4, 29. , 16.4, 25. ,
        25. , 17.1, 23.2, 10.4, 19.6, 17.2, 27.5, 23. , 50. , 17.9,  9.6,
        17.2, 22.5, 21.4])
```

```
y_pred = model.predict(X_test)
```

```
y_pred
```

```
array([28.99672362, 36.02556534, 14.81694405, 25.03197915, 18.76987992,
        23.25442929, 17.66253818, 14.34119 , 23.01320703, 20.63245597,
        24.90850512, 18.63883645, -6.08842184, 21.75834668, 19.23922576,
        26.19319733, 20.64773313,  5.79472718, 40.50033966, 17.61289074,
        27.24909479, 30.06625441, 11.34179277, 24.16077616, 17.86058499,
        15.83609765, 22.78148106, 14.57704449, 22.43626052, 19.19631835,
        22.43383455, 25.21979081, 25.93909562, 17.70162434, 16.76911711,
        16.95125411, 31.23340153, 20.13246729, 23.76579011, 24.6322925 ,
        13.94204955, 32.25576301, 42.67251161, 17.32745046, 27.27618614,
        16.99310991, 14.07009109, 25.90341861, 20.29485982, 29.95339638,
        21.28860173, 34.34451856, 16.04739105, 26.22562412, 39.53939798,
        22.57950697, 18.84531367, 32.72531661, 25.0673037 , 12.88628956,
        22.68221908, 30.48287757, 31.52626806, 15.90148607, 20.22094826,
        16.71089812, 20.52384893, 25.96356264, 30.61607978, 11.59783023,
        20.51232627, 27.48111878, 11.01962332, 15.68096344, 23.79316251,
         6.19929359, 21.6039073 , 41.41377225, 18.76548695,  8.87931901,
        20.83076916, 13.25620627, 20.73963699,  9.36482222, 23.22444271,
        31.9155003 , 19.10228271, 25.51579303, 29.04256769, 20.14358566,
        25.5859787 ,  5.70159447, 20.09474756, 14.95069156, 12.50395648,
        20.72635294, 24.73957161, -0.164237 , 13.68486682, 16.18359697,
        22.27621999, 24.47902364])
```

## ✓ Testing the model performance

```
model.score(X_test,y_test)
```

```
0.6687594935356307
```

```
# R squared
```

```
r2_score(y_test,y_pred)
```

```
0.6687594935356307
```

```
# Adjusted R squared
```

```
# MSE
```

```
mean_squared_error(y_test,y_pred)
```

```
24.291119474973616
```

```
# MAE
```

```
mean_absolute_error(y_test,y_pred)
```

```
3.1890919658878523
```

```
# RMSE
```

```
np.sqrt(mean_squared_error(y_test,y_pred))
```

```
4.9286021826653466
```

```
plt.scatter(y_test,y_pred)
```

```
plt.xlabel("Actual Price")
```

```
plt.ylabel("Predicted Price")
```

```
plt.grid()
```

```
plt.plot([min(y_test),max(y_test)], [min(y_pred),max(y_pred)], color='red')
```

```
plt.title('Actual Price V/s Predicted Price')
```

```
Text(0.5, 1.0, 'Actual Price V/s Predicted Price')
```

